


Contents

- Comparison of CPUs and GPUs
- Programming Styles
- What is Easy to Accelerate?
- Libraries


CPU or GPU?

Which die is the CPU, which one the GPU?

GK110




XEON-E7




Four Levels of Parallelism


SIMD 10x



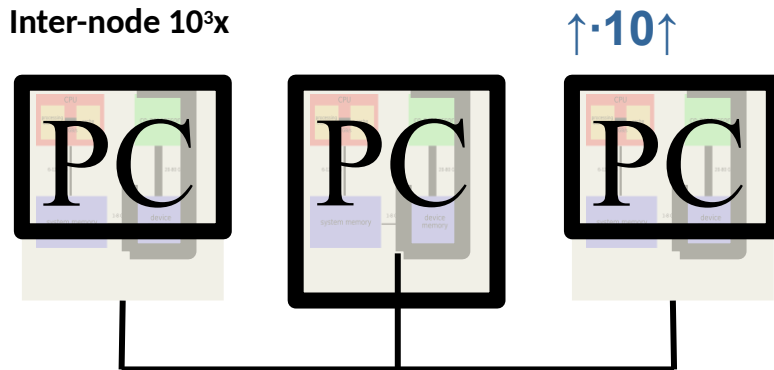
Many-core 10x




Intra-node 10x



Inter-node 10³x



Bandwidth in an Accelerator System




GPUs vs. CPUs


	Tesla K20	Xeon E7-4800 4P
Core count	13 SMs 64/832 (DP), 192/2,496 (SP)	10 Cores 2 FP-ALUs/core, SSE 16B
Frequency	0.7GHz	2.4GHz
Peak Compute Performance	1,165 GFLOPS (DP) 3,494 GFLOPS (SP)	96 GFLOPS (DP)
Use model	throughput-oriented	latency-oriented
Latency treatment	toleration	minimization
Programming	1000s-10,000s of threads	10s of threads
Memory bandwidth	250 GBytes/sec	34 GByte/s (per P)
Memory capacity	5 GB	up to 2TB
Die size	550mm ²	684 mm ²
Transistor count	7.1 billion	2.3 billion
Technology	28nm	32nm
Power consumption	250W	130W
Power efficiency	4.66 GFLOPs/Watt (DP) 14 GFLOPs/Watt (SP)	0.74 GFLOPs/Watt (DP)

Theoretical Performance

Theoretical GFLOP/s



Theoretical GB/s



Parallelism on CPUs and GPUs

CPU

- SIMD AVX 32B, Phi 64B
 - MADD with 8-16 floats MADD
 - with 4-8 doubles Coding: **explicit**,
 - automatic
- **Minimum** multi-threading
 - $\#threads = \#cores$ (good)
 - $\#threads = 2 * \#cores$ (good)
 - $\#threads = 10 * \#cores$ (difficult)
 - Coding: **explicit with resident threads**, implicit with libraries

GPU

- SIMD warp size 32
 - MADD with 32 floats
 - MADD with 32 doubles
 - Coding: **implicit**, partly explicit
- **Maximum** multi-threading
 - $\#warps = \#cores \approx 15$ (bad)
 - $\#warps \approx 100$ (difficult)
 - $\#warps > 1000$ (good)
 - Coding: **implicit with max. parallelism**, explicit (advanced)

Memory on CPUs and GPUs

CPU

- Deep and large memories
 - Core: Reg, L1, L2
 - Shared: L3, eDRAM
 - Coding: **implicit**
- Usage
 - Optimize 1D locality
 - Optimize size of working sets
 - Prefetch, pipeline
 - Coding: implicit, explicit

GPU

- Smaller and specialized memories
 - Core: Reg, L1 or shmem
 - Shared: L2, constant
 - Coding: **explicit**, implicit with libs
- Usage
 - Optimize 1D, 2D, 3D locality
 - Decide on data location: Reg, L1, shmem, constant
 - Many warps and low latency vs. amount of local data per warp
 - Coding: explicit

Similarities and Differences

CPU

- For high performance
 - SIMD
 - Multi-threading
 - Memory access alignment Minimal
 - latency with **large caches** Working set
 - opt. wrt **deep caches**
 - Locality opt. wrt **cache lines, NUMA**
 -
- Opt. **serial** performance
 - High normal and boost frequency
 - Low latency caches
 - Speculative execution

GPU

- For high performance
 - SIMD
 - Multi-threading
 - Memory access alignment
 - Minimal latency with **many warps**
 - **Working set opt. wrt #warps**
 - **Locality opt. wrt memory types**
 - Opt. **throughput** performance
- Lower normal and boost frequency
- L2 latency is high
- No speculative execution

CUDA Ecosystem

What is Easy to Accelerate?

Embarrassingly Parallel Loop

- ```
for(int i=0; i<SIZE; ++i)
{
 c[i]= a[i+1]+b[i]*a[i-yoff];
 func(a,b,c,i);
}
```
- Relevant for performance
  - `SIZE > 10k`
  - Arithmetic intensity
  - Regularity of memory access
  - Amount of local state and reuse

# Branches in Loops

- ```
for(int i=0; i<SIZE; ++i) {  
    if(cond(i)) special_func(a,b,c,i);  
    else normal_func(a,b,c,i);  
}
```
- **Only a problem if all these conditions hold:**
 - Special case is more than 10% of cases
 - Normal and special case differ largely in execution times
 - Data of special cases is scattered in memory

Index Dependencies

- ```
for (int i=1; i<SIZE; ++i)
{ a[i] += a[i-1];
}
```
- Replace serial dependence
  - Use equivalent parallel variant
  - If allowed, use approximate parallel variant
  - Check if parent computation can use other ingredients

# Data Movement is Critical

- ```
CPU_func1(a,b,c);  
GPU_func1(a,b,c); // implicit transfer!  
CPU_func2(a,b,c);  
GPU_func2(a,b,c); // implicit transfer!
```
- Such CPU-GPU alternation only works well if
 - Execution time of `GPU_func*` is at least a millisecond
 - High arithmetic intensity, e.g. `matrix*matrix`, not `matrix*vector`
- Otherwise
 - GPU must perform multiple operations on the same data, e.g. multiple vector-vector or matrix-vector operations.

Linked List

- ```
for (; elm!=nullptr; elm= elm->next) {
 func(elm->data);
}
```
- **Do not do this!**
  - Unless all parallelism can be used efficiently in `func()`
  - Terrible performance on CPUs and GPUs
  - Vector almost always dramatically faster than list
  - Even `insert(pos)`, `delete(pos)` much faster in vector if we first search for `pos`



# Data Structures on GPUs


|                      | Difficulty | Speed | Support       | Format                                  |
|----------------------|------------|-------|---------------|-----------------------------------------|
| <b>vector</b>        | easy       | fast  | everywhere    | contiguous data                         |
| <b>dense matrix</b>  | easy       | fast  | many libs     | contiguous data                         |
| <b>sparse matrix</b> | moderate   | fast  | many libs     | <b>CSR</b> , BCSR, (B)CSC, COO, special |
| <b>graph</b>         | moderate   | fast  | multiple libs | <b>CSR</b> , special                    |
| <b>tree</b>          | difficult  | fast  | little        | various special formats                 |
| <b>list</b>          | moderate   | slow  | none          | special formats                         |

# Where to Put the Parallelism?


```
• for(int i=0; i<SIZE; i+=block_size)
{ func1(a,b,c,i,block_size);
 for(int i=0; i<SIZE; i+=block_size)
{ func2(a,b,c,i,block_size);
 }
 for(int i=0; i<SIZE; i+=block_size) {

func3(a,b,c,i,block_size);
 }
}
```


3 loops, block\_size=1



3 loops, block\_size=10



1 loop, block\_size=10



# Libraries

# Summary

- For high performance on CPUs and GPUs
  - High parallelism and high data locality
  - Optimizations are similar in concept, but different and very involved in detail
  - Difficult to do by hand → use libraries
- What is easy to accelerate?
  - Large loops with no/simple index dependencies
  - Data placement and movement are crucial